

Word2Vec with NLTK Reuters Corpus and Visualization using t-SNE

Lab Activity

Import Required Libraries

```
import nltk
from nltk.corpus import reuters, stopwords
from gensim.models import Word2Vec
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import pandas as pd
```

Download the Reuters corpus (if not already downloaded)

```
nltk.download('reuters')
nltk.download('punkt')
nltk.download('punkt_tab')
nltk.download('stopwords')
```

Step 1: Load the Reuters Corpus

Extract sentences from the Reuters corpus Explanation: The Reuters corpus contains a collection of financial news articles. Each sentence will be tokenized into words to create the vocabulary for Word2Vec.

```
corpus_sentences = []
for fileid in reuters.fileids():
    raw_text = reuters.raw(fileid)
    tokenized_sentence = [word for word in nltk.word_tokenize(raw_text) if word.isalnum() and word]
    corpus_sentences.append(tokenized_sentence)
print(f"Number of sentences in the Reuters corpus: {len(corpus_sentences)})")
```

Step 2: Train a Word2Vec Model

Explanation: Word2Vec will learn vector representations for each word in the vocabulary.

Parameters:

- vector_size: Size of the word embedding vectors.
- window: Context window size for training.
- min_count: Minimum frequency for a word to be included in the vocabulary.
- workers: Number of threads for parallel processing.

```
model = Word2Vec(sentences=corpus_sentences, vector_size=100, window=5, min_count=5, workers=4)
# Print vocabulary size
print(f"Vocabulary size: {len(model.wv.index_to_key)}")
```

Step 3: Extract Word Embeddings for Visualization

```
import numpy as np
# Extract the learned word vectors and their corresponding words for visualization.
words = list(model.wv.index_to_key)[:200] # Limit to top 200 words for better visualization
word_vectors = np.array([model.wv[word] for word in words]) # Convert to NumPy array for compatibility
```

Step 4: Reduce Dimensionality with t-SNE

```
# Use t-SNE to project the high-dimensional word embeddings into a 2D space.
tsne = TSNE(n_components=2, random_state=42, perplexity=30)
word_vectors_2d = tsne.fit_transform(word_vectors)
```

Step 5: Visualize the Word Embeddings

```
# Plot the 2D t-SNE visualization of the word embeddings with their labels.
def plot_embeddings(vectors, labels):
    plt.figure(figsize=(16, 12))
    for i, label in enumerate(labels):
        x, y = vectors[i]
        plt.scatter(x, y, color='blue')
        plt.text(x + 0.1, y + 0.1, label, fontsize=9)
    plt.title("Word2Vec Embeddings Visualized with t-SNE")
    plt.xlabel("t-SNE Dimension 1")
    plt.ylabel("t-SNE Dimension 2")
    plt.show()

plot_embeddings(word_vectors_2d, words)
```

Notes:

1. The Word2Vec model captures semantic relationships between words based on their co-occurrence in the text.
2. Similar words tend to cluster together in the t-SNE visualization.
3. t-SNE helps reduce the dimensionality for better interpretability, but the visualization may slightly vary with different runs due to its stochastic nature.

Key Questions for Students:

- What do you observe about the clusters in the t-SNE plot?
- How do you think the choice of parameters (e.g., window size, vector size) affects the embeddings?
- What are the limitations of using Word2Vec and t-SNE for NLP tasks?

Information Retrieval Task:

Task: Build a Document Retrieval System using Word2Vec

1. Given a query string, find the most relevant documents from the Reuters corpus using Word2Vec embeddings.
2. Steps:
 - a. Preprocess the query string by tokenizing and removing stop words.
 - b. Compute the average Word2Vec embedding for the query string.
 - c. Compute the average Word2Vec embedding for each document in the Reuters corpus.
 - d. Use cosine similarity to find the top N most relevant documents for the query.
3. Display the top N document IDs and their similarity scores.

